

Aprendizaje reforzado con Super Mario






José Carlos Batista Rivero

Índice

- Introducción
- Aprendizaje reforzado
- DQN (Deep Q Learning)
- Demo

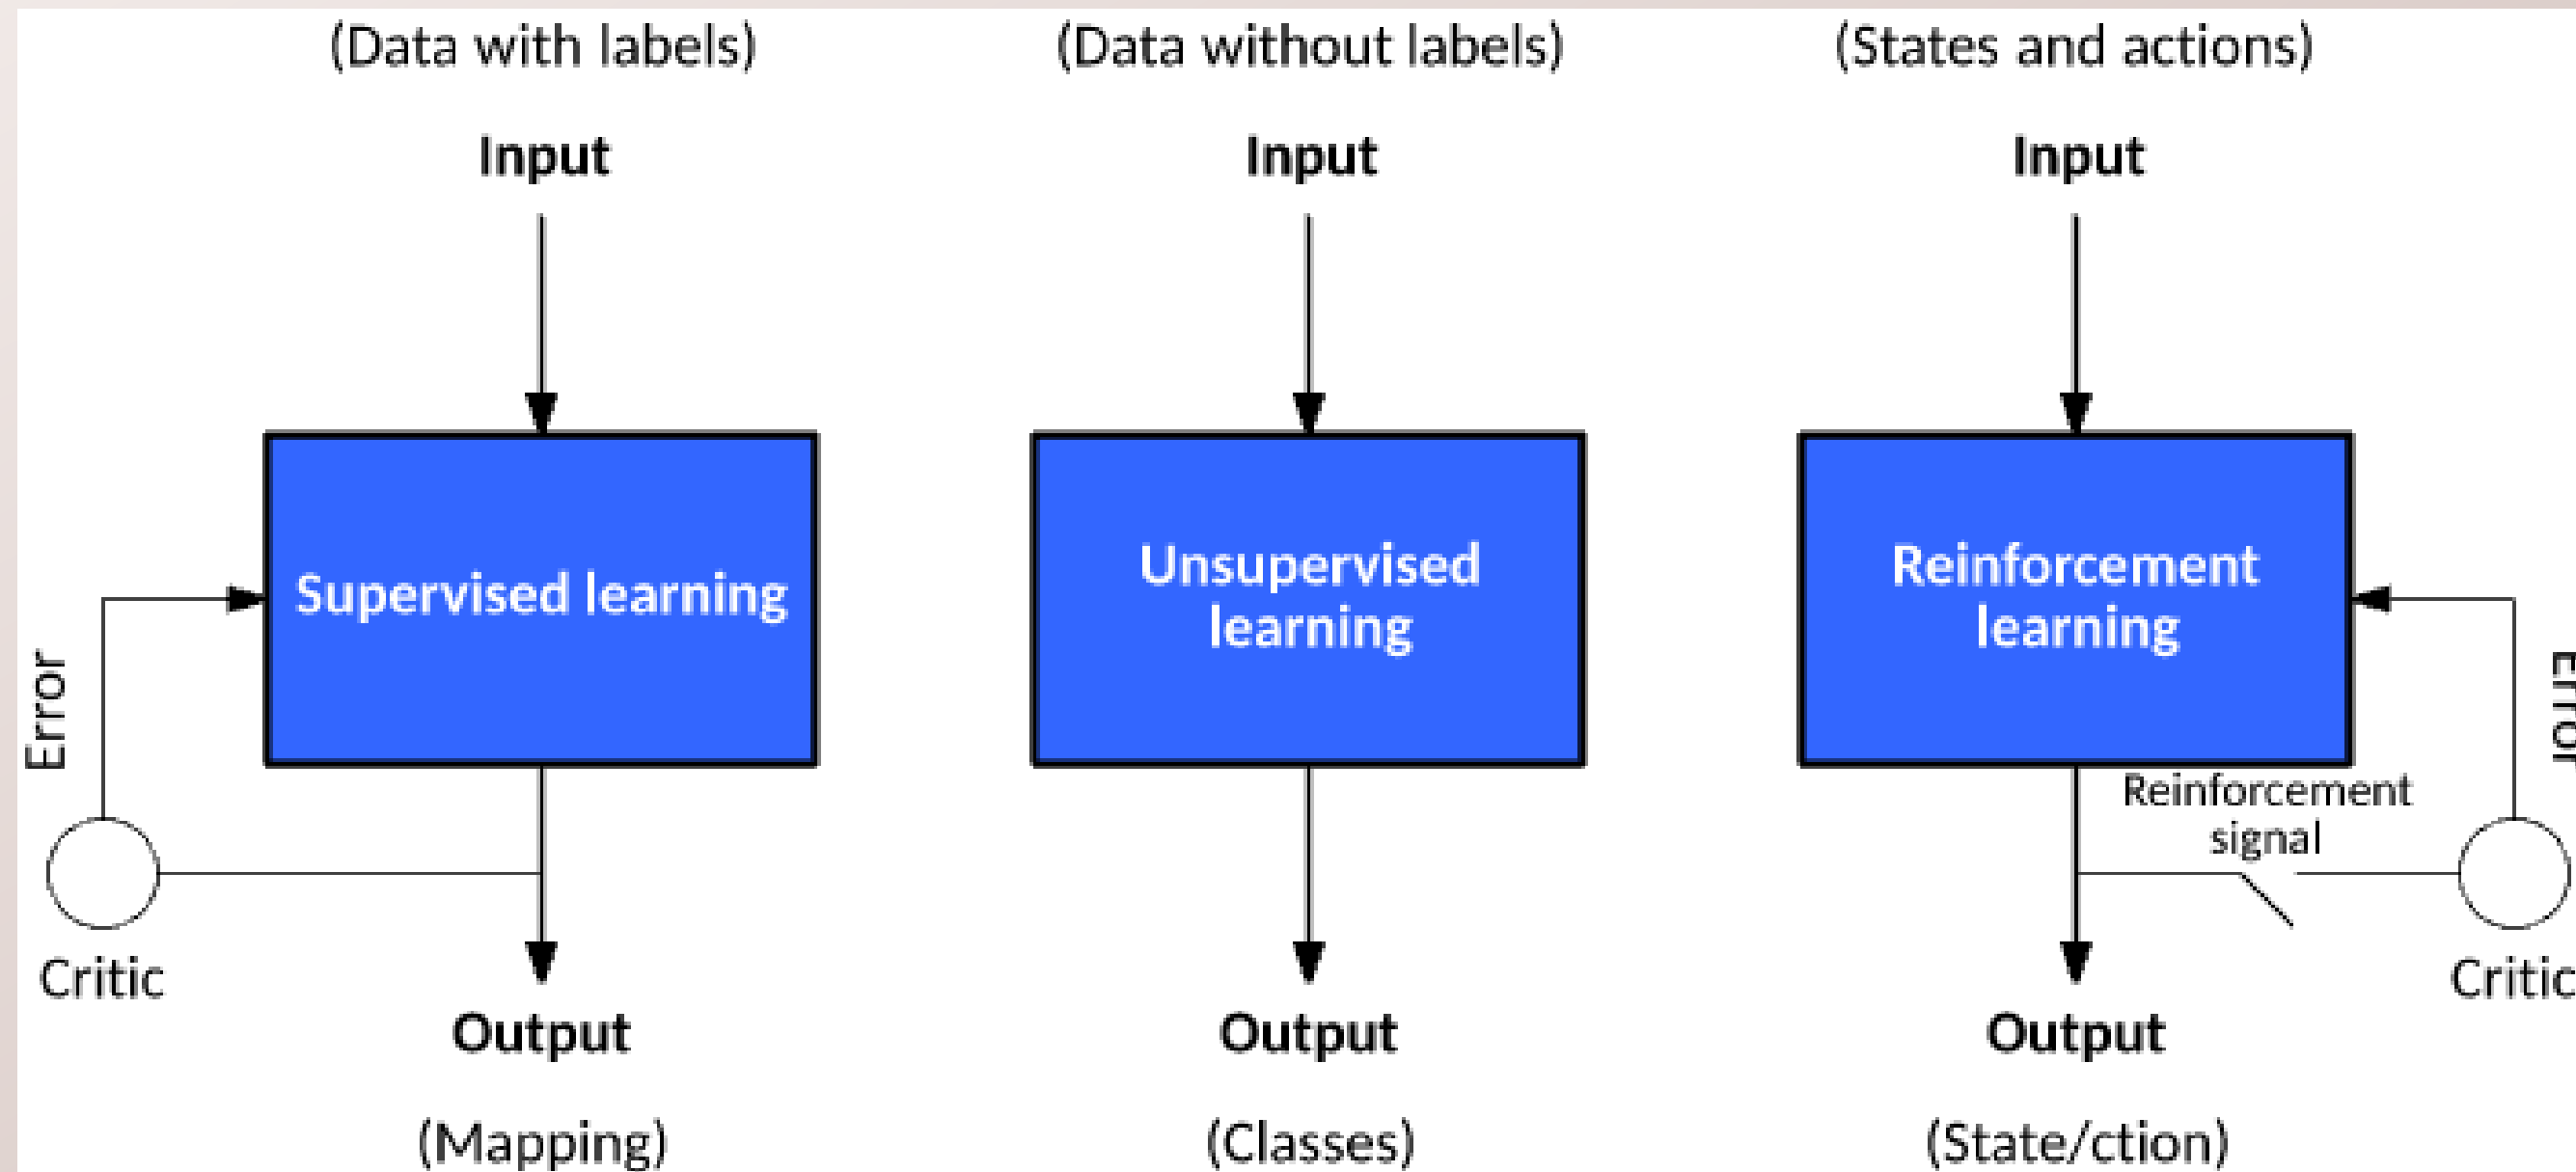
Introducción

OpenAI Gym

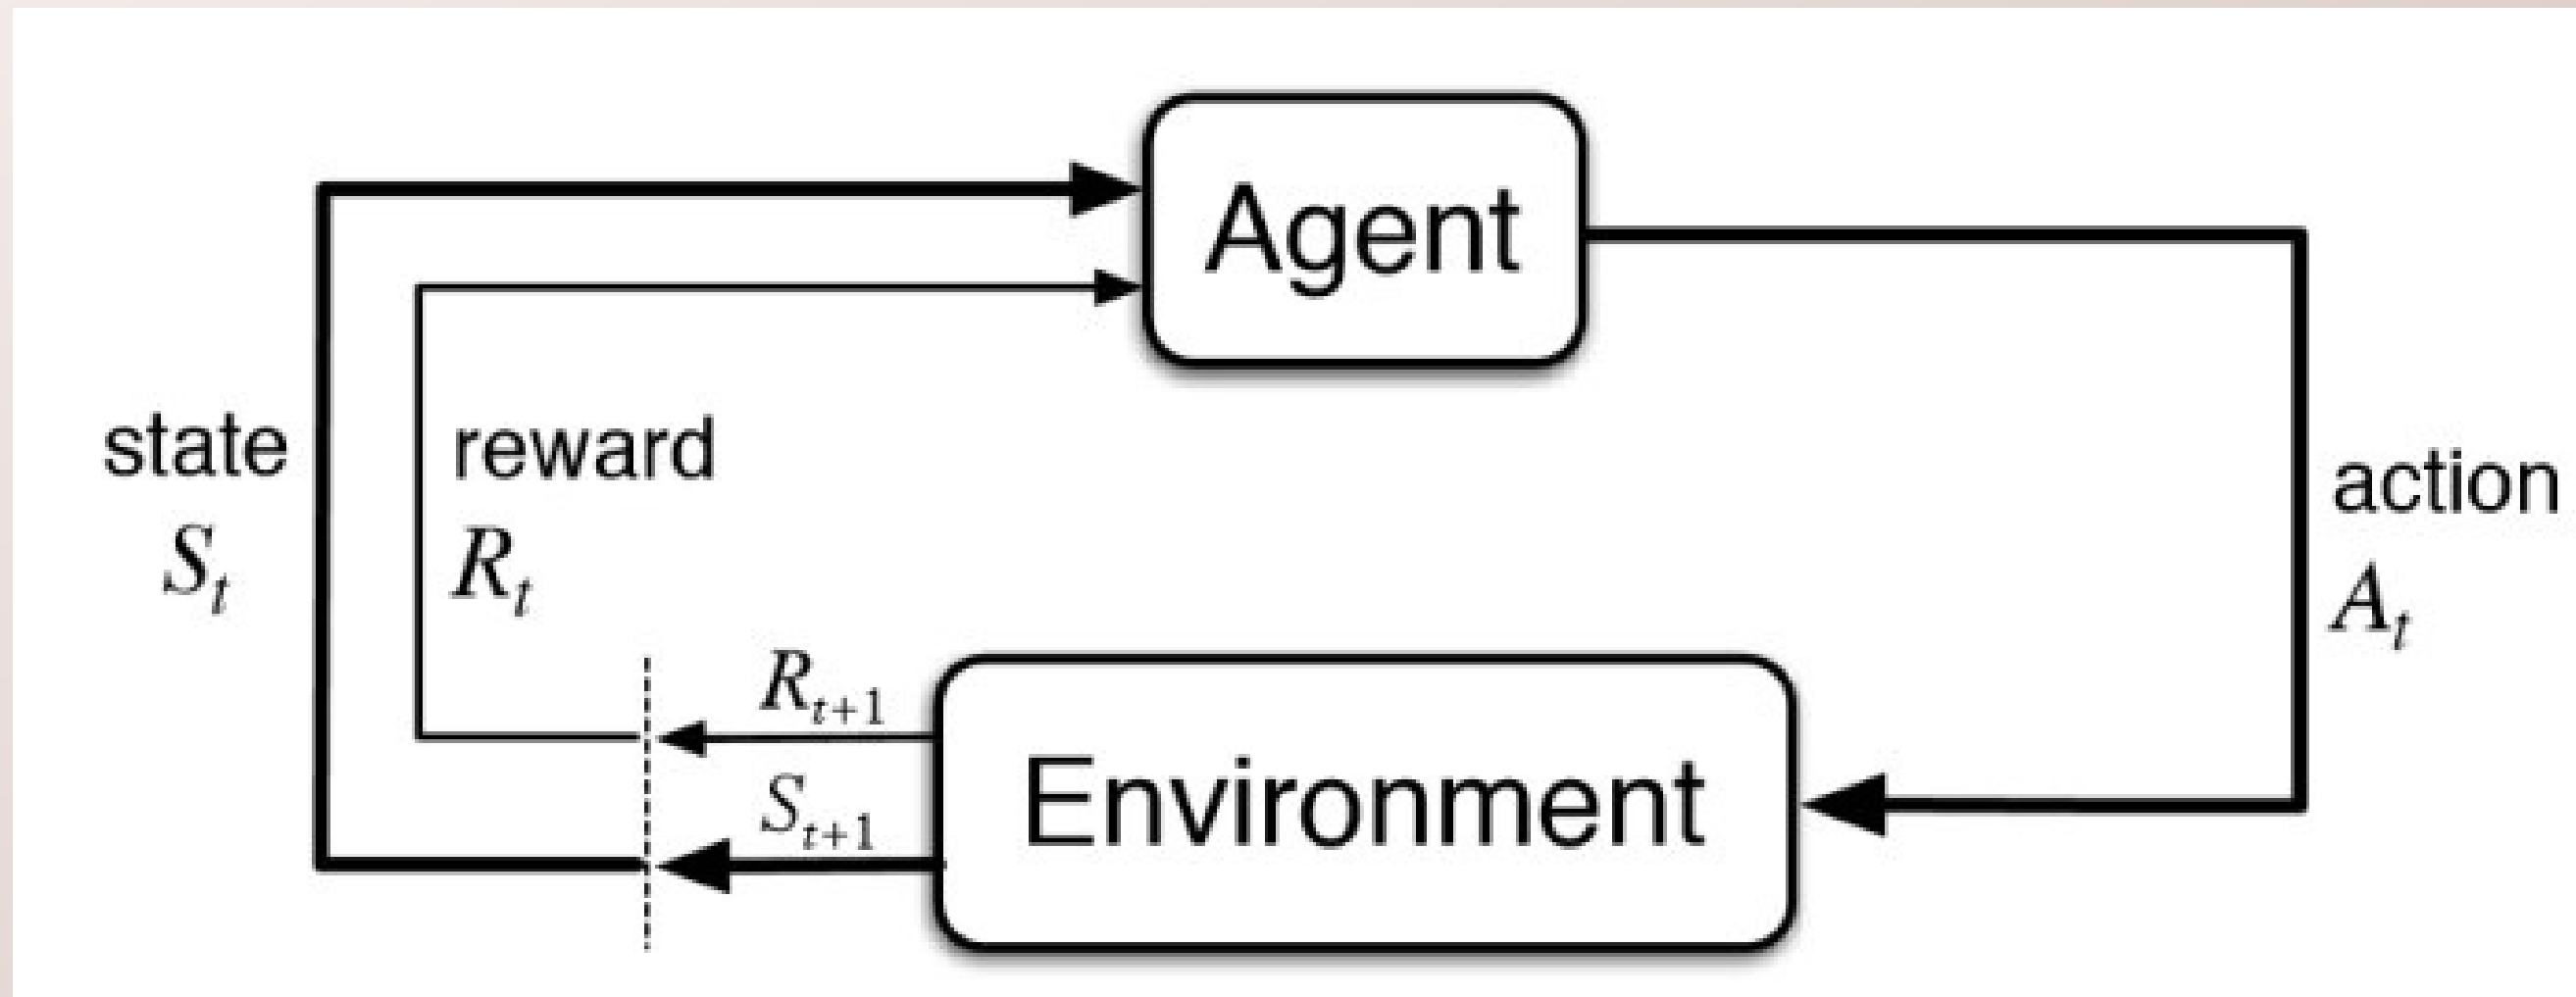


Aprendizaje reforzado

¿Qué es el aprendizaje reforzado?




Filosofía



Agente



Entorno




Acción




Recompensa

$$r = x + t + d$$



Recompensa



Deep Q Learning

Valor Q

$$Q^{\pi}(s_t, a_t) = \underline{E}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots | s_t, a_t]$$

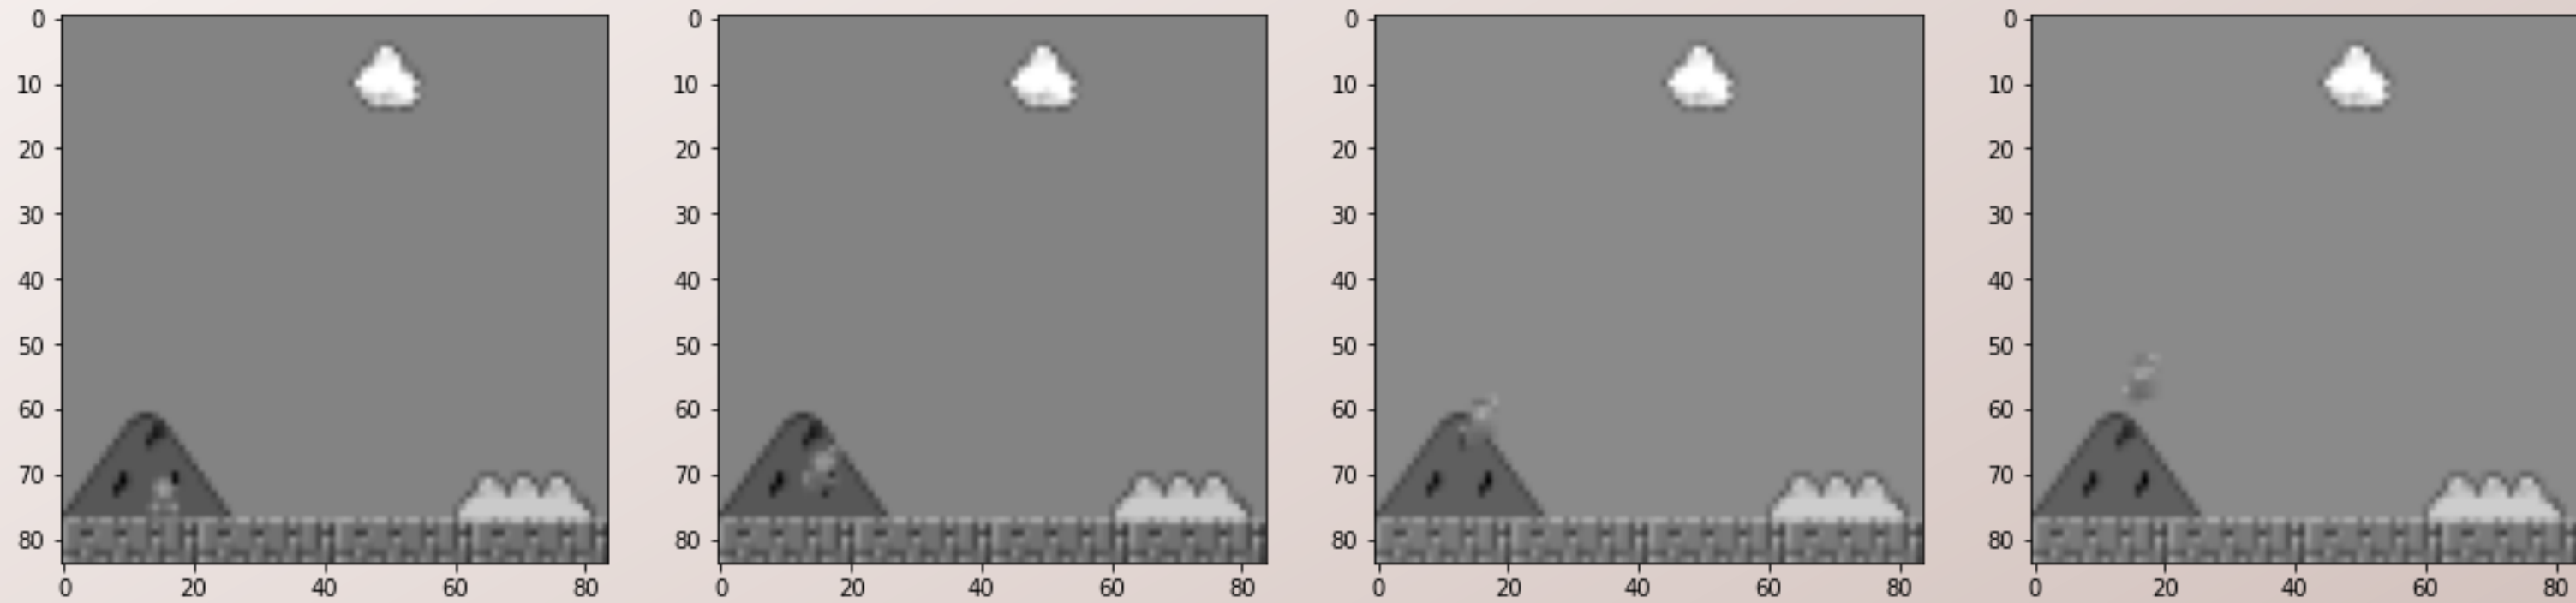
Q-Values for the state
given a particular state

Expected discounted
cumulative reward

Given the state and action

$$0 < \gamma < 1$$

¿Por qué una red neuronal?




$$256^{4 \cdot 84 \cdot 84} = 256^{28224}$$

S

$F(S)$

A



Pytorch vs Keras

Keras

```
def build_model(input_shape, n_actions):  
    model = keras.Sequential()  
    model.add(keras.layers.Conv2D(32, kernel_size=32, strides=4, activation='ReLU', input_shape=input_shape))  
    model.add(keras.layers.Conv2D(32, kernel_size=64, strides=2, activation='ReLU'))  
    model.add(keras.layers.Conv2D(32, kernel_size=64, strides=1, activation='ReLU'))  
    model.add(keras.layers.Flatten())  
    model.add(keras.layers.Dense(512, activation='ReLU'))  
    model.add(keras.layers.Conv2D(n_actions, activation='linear'))  
  
    return model
```

Pytorch

```
class ConvNet(nn.Module):  
    def __init__(self, input_shape, n_actions):  
        super(ConvNet, self).__init__()  
        self.conv = nn.Sequential(  
            nn.Conv2d(input_shape[0], 32, kernel_size=8, stride=4),  
            nn.ReLU(),  
            nn.Conv2d(32, 64, kernel_size=4, stride=2),  
            nn.ReLU(),  
            nn.Conv2d(64, 64, kernel_size=3, stride=1),  
            nn.ReLU()  
        )  
  
        conv_out_size = self._get_conv_out(input_shape)  
        self.fc = nn.Sequential(  
            nn.Linear(conv_out_size, 512),  
            nn.ReLU(),  
            nn.Linear(512, n_actions)  
        )  
  
    def _get_conv_out(self, shape):  
        o = self.conv(torch.zeros(1, *shape))  
        return int(np.prod(o.size()))  
  
    def forward(self, x):  
        conv_out = self.conv(x).view(x.size()[0], -1)  
        return self.fc(conv_out)
```

Función de coste

$$loss = \left(\underbrace{r + \gamma \max_{a'} \hat{Q}(s, a')}_{\text{Target}} - \underbrace{Q(s, a)}_{\text{Prediction}} \right)^2$$

Diagram illustrating the cost function (loss) for a reinforcement learning agent. The formula is:

$$loss = \left(\underbrace{r + \gamma \max_{a'} \hat{Q}(s, a')}_{\text{Target}} - \underbrace{Q(s, a)}_{\text{Prediction}} \right)^2$$

The components are labeled:

- Reward**: r
- Decay Rate**: γ
- Target**: $r + \gamma \max_{a'} \hat{Q}(s, a')$
- Prediction**: $Q(s, a)$


Mejorar el entrenamiento

- Exploración vs Predicción
- Memory replay
- Doble red neuronal

Parámetros

	Modelo	learning_rate	gamma	epsilon_decay	score_episode
0	Final	0.00050	0.95	0.999	21.974
1	Prueba1	0.10000	0.92	0.995	-38.436
2	Prueba2	0.00100	0.97	0.990	19.740
3	Prueba3	0.00010	0.99	0.950	21.736
4	Prueba4	0.00001	1.00	0.900	17.922

Entrenamiento



Demo

*"Lo siento Mario, pero la
princesa está en otro castillo"*

Un Toad

¡Muchas gracias!

Autor : José Carlos Batista Rivero

Correo electrónico : josecarlosbatistarivero@gmail.com

github : <https://github.com/JoseKuru>